

O'REILLY®

Free Sampler

Elasticsearch

The Definitive Guide

A DISTRIBUTED REAL-TIME SEARCH AND ANALYTICS ENGINE

Clinton Gormley &
Zachary Tong

O'Reilly Ebooks—Your bookshelf on your devices!



When you buy an ebook through oreilly.com you get lifetime access to the book, and whenever possible we provide it to you in five, DRM-free file formats—PDF, .epub, Kindle-compatible .mobi, Android .apk, and DAISY—that you can use on the devices of your choice. Our ebook files are fully searchable, and you can cut-and-paste and print them. We also alert you when we've updated the files with corrections and additions.

Learn more at ebooks.oreilly.com

You can also purchase O'Reilly ebooks through the iBookstore, the [Android Marketplace](http://AndroidMarketplace.com), and Amazon.com.

O'REILLY®

Spreading the knowledge of innovators

oreilly.com

Elasticsearch: The Definitive Guide

by Clinton Gormley and Zachary Tong

Copyright © 2015 Elasticsearch. All rights reserved.

Printed in the United States of America.

Published by O'Reilly Media, Inc. , 1005 Gravenstein Highway North, Sebastopol, CA 95472.

O'Reilly books may be purchased for educational, business, or sales promotional use. Online editions are also available for most titles (<http://safaribooksonline.com>). For more information, contact our corporate/institutional sales department: 800-998-9938 or corporate@oreilly.com.

Editors: Mike Loukides and Brian Anderson

Production Editor: Shiny Kalapurakkal

Proofreader: Sharon Wilkey

Indexer: Ellen Troutman-Zaig

Interior Designer: David Futato

Cover Designer: Ellie Volkhausen

Illustrator: Rebecca Demarest

January 2015: First Edition

Revision History for the First Edition

2015-01-16: First Release

See <http://oreilly.com/catalog/errata.csp?isbn=9781449358549> for release details.

The O'Reilly logo is a registered trademark of O'Reilly Media, Inc. *Elasticsearch: The Definitive Guide*, the cover image, and related trade dress are trademarks of O'Reilly Media, Inc.

Many of the designations used by manufacturers and sellers to distinguish their products are claimed as trademarks. Where those designations appear in this book, and O'Reilly Media, Inc. was aware of a trademark claim, the designations have been printed in caps or initial caps.

While the publisher and the authors have used good faith efforts to ensure that the information and instructions contained in this work are accurate, the publisher and the authors disclaim all responsibility for errors or omissions, including without limitation responsibility for damages resulting from the use of or reliance on this work. Use of the information and instructions contained in this work is at your own risk. If any code samples or other technology this work contains or describes is subject to open source licenses or the intellectual property rights of others, it is your responsibility to ensure that your use thereof complies with such licenses and/or rights.

978-1-449-35854-9

[LSI]

Table of Contents

Foreword.....	xxi
---------------	-----

Preface.....	xxiii
--------------	-------

Part I. Getting Started

1. You Know, for Search.....	3
Installing Elasticsearch	4
Installing Marvel	5
Running Elasticsearch	5
Viewing Marvel and Sense	6
Talking to Elasticsearch	6
Java API	6
RESTful API with JSON over HTTP	7
Document Oriented	9
JSON	9
Finding Your Feet	10
Let's Build an Employee Directory	10
Indexing Employee Documents	10
Retrieving a Document	12
Search Lite	13
Search with Query DSL	15
More-Complicated Searches	16
Full-Text Search	17
Phrase Search	18
Highlighting Our Searches	19
Analytics	20
Tutorial Conclusion	23

Distributed Nature	23
Next Steps	24
2. Life Inside a Cluster.....	25
An Empty Cluster	26
Cluster Health	26
Add an Index	27
Add Failover	29
Scale Horizontally	30
Then Scale Some More	31
Coping with Failure	32
3. Data In, Data Out.....	35
What Is a Document?	36
Document Metadata	37
_index	37
_type	37
_id	38
Other Metadata	38
Indexing a Document	38
Using Our Own ID	38
Autogenerating IDs	39
Retrieving a Document	40
Retrieving Part of a Document	41
Checking Whether a Document Exists	42
Updating a Whole Document	42
Creating a New Document	43
Deleting a Document	44
Dealing with Conflicts	45
Optimistic Concurrency Control	47
Using Versions from an External System	49
Partial Updates to Documents	50
Using Scripts to Make Partial Updates	51
Updating a Document That May Not Yet Exist	52
Updates and Conflicts	53
Retrieving Multiple Documents	54
Cheaper in Bulk	56
Don't Repeat Yourself	60
How Big Is Too Big?	60
4. Distributed Document Store.....	61
Routing a Document to a Shard	61

How Primary and Replica Shards Interact	62
Creating, Indexing, and Deleting a Document	63
Retrieving a Document	65
Partial Updates to a Document	66
Multidocument Patterns	67
Why the Funny Format?	69
5. Searching—The Basic Tools.....	71
The Empty Search	72
hits	73
took	73
shards	73
timeout	74
Multi-index, Multitype	74
Pagination	75
Search <i>Lite</i>	76
The <code>_all</code> Field	77
More Complicated Queries	78
6. Mapping and Analysis.....	79
Exact Values Versus Full Text	80
Inverted Index	81
Analysis and Analyzers	84
Built-in Analyzers	84
When Analyzers Are Used	85
Testing Analyzers	86
Specifying Analyzers	87
Mapping	87
Core Simple Field Types	88
Viewing the Mapping	89
Customizing Field Mappings	89
Updating a Mapping	91
Testing the Mapping	92
Complex Core Field Types	93
Multivalue Fields	93
Empty Fields	93
Multilevel Objects	94
Mapping for Inner Objects	94
How Inner Objects are Indexed	95
Arrays of Inner Objects	95

7. Full-Body Search.....	97
Empty Search	97
Query DSL	98
Structure of a Query Clause	99
Combining Multiple Clauses	99
Queries and Filters	100
Performance Differences	101
When to Use Which	101
Most Important Queries and Filters	102
term Filter	102
terms Filter	102
range Filter	102
exists and missing Filters	103
bool Filter	103
match_all Query	103
match Query	104
multi_match Query	104
bool Query	105
Combining Queries with Filters	105
Filtering a Query	106
Just a Filter	107
A Query as a Filter	107
Validating Queries	108
Understanding Errors	108
Understanding Queries	109
 8. Sorting and Relevance.....	 111
Sorting	111
Sorting by Field Values	112
Multilevel Sorting	113
Sorting on Multivalue Fields	113
String Sorting and Multifields	114
What Is Relevance?	115
Understanding the Score	116
Understanding Why a Document Matched	119
Fielddata	119
 9. Distributed Search Execution.....	 121
Query Phase	122
Fetch Phase	123
Search Options	125
preference	125

timeout	126
routing	126
search_type	127
scan and scroll	127
10. Index Management.....	131
Creating an Index	131
Deleting an Index	132
Index Settings	132
Configuring Analyzers	133
Custom Analyzers	134
Creating a Custom Analyzer	135
Types and Mappings	137
How Lucene Sees Documents	137
How Types Are Implemented	138
Avoiding Type Gotchas	138
The Root Object	140
Properties	140
Metadata: _source Field	141
Metadata: _all Field	142
Metadata: Document Identity	144
Dynamic Mapping	145
Customizing Dynamic Mapping	147
date_detection	147
dynamic_templates	148
Default Mapping	149
Reindexing Your Data	150
Index Aliases and Zero Downtime	151
11. Inside a Shard.....	153
Making Text Searchable	154
Immutability	155
Dynamically Updatable Indices	155
Deletes and Updates	158
Near Real-Time Search	159
refresh API	160
Making Changes Persistent	161
flush API	165
Segment Merging	166

Part II. Search in Depth

12. Structured Search.....	173
Finding Exact Values	173
term Filter with Numbers	174
term Filter with Text	175
Internal Filter Operation	178
Combining Filters	179
Bool Filter	179
Nesting Boolean Filters	181
Finding Multiple Exact Values	182
Contains, but Does Not Equal	183
Equals Exactly	184
Ranges	185
Ranges on Dates	186
Ranges on Strings	187
Dealing with Null Values	187
exists Filter	188
missing Filter	190
exists/missing on Objects	191
All About Caching	192
Independent Filter Caching	192
Controlling Caching	193
Filter Order	194
13. Full-Text Search.....	197
Term-Based Versus Full-Text	197
The match Query	199
Index Some Data	199
A Single-Word Query	200
Multiword Queries	201
Improving Precision	202
Controlling Precision	203
Combining Queries	204
Score Calculation	205
Controlling Precision	205
How match Uses bool	206
Boosting Query Clauses	207
Controlling Analysis	209

Default Analyzers	211
Configuring Analyzers in Practice	213
Relevance Is Broken!	214
14. Multifield Search.....	217
Multiple Query Strings	217
Prioritizing Clauses	218
Single Query String	219
Know Your Data	220
Best Fields	221
dis_max Query	222
Tuning Best Fields Queries	223
tie_breaker	224
multi_match Query	225
Using Wildcards in Field Names	226
Boosting Individual Fields	227
Most Fields	227
Multifield Mapping	228
Cross-fields Entity Search	231
A Naive Approach	231
Problems with the most_fields Approach	232
Field-Centric Queries	232
Problem 1: Matching the Same Word in Multiple Fields	233
Problem 2: Trimming the Long Tail	233
Problem 3: Term Frequencies	234
Solution	235
Custom _all Fields	235
cross-fields Queries	236
Per-Field Boosting	238
Exact-Value Fields	239
15. Proximity Matching.....	241
Phrase Matching	242
Term Positions	242
What Is a Phrase	243
Mixing It Up	244
Multivalued Fields	245
Closer Is Better	246
Proximity for Relevance	247
Improving Performance	249
Rescoring Results	249
Finding Associated Words	250

Producing Shingles	251
Multifields	252
Searching for Shingles	253
Performance	255
16. Partial Matching.....	257
Postcodes and Structured Data	258
prefix Query	259
wildcard and regexp Queries	260
Query-Time Search-as-You-Type	262
Index-Time Optimizations	264
Ngrams for Partial Matching	264
Index-Time Search-as-You-Type	265
Preparing the Index	265
Querying the Field	267
Edge n-grams and Postcodes	270
Ngrams for Compound Words	271
17. Controlling Relevance.....	275
Theory Behind Relevance Scoring	275
Boolean Model	276
Term Frequency/Inverse Document Frequency (TF/IDF)	276
Vector Space Model	279
Lucene's Practical Scoring Function	282
Query Normalization Factor	283
Query Coordination	284
Index-Time Field-Level Boosting	286
Query-Time Boosting	286
Boosting an Index	287
t.getBoost()	288
Manipulating Relevance with Query Structure	288
Not Quite Not	289
boosting Query	290
Ignoring TF/IDF	291
constant_score Query	291
function_score Query	293
Boosting by Popularity	294
modifier	296
factor	298
boost_mode	299
max_boost	301
Boosting Filtered Subsets	301

filter Versus query	302
functions	303
score_mode	303
Random Scoring	303
The Closer, The Better	305
Understanding the price Clause	308
Scoring with Scripts	308
Pluggable Similarity Algorithms	310
Okapi BM25	310
Changing Similarities	313
Configuring BM25	314
Relevance Tuning Is the Last 10%	315

Part III. Dealing with Human Language

18. Getting Started with Languages.....	319
Using Language Analyzers	320
Configuring Language Analyzers	321
Pitfalls of Mixing Languages	323
At Index Time	323
At Query Time	324
Identifying Language	324
One Language per Document	325
Foreign Words	326
One Language per Field	327
Mixed-Language Fields	329
Split into Separate Fields	329
Analyze Multiple Times	329
Use n-grams	330
19. Identifying Words.....	333
standard Analyzer	333
standard Tokenizer	334
Installing the ICU Plug-in	335
icu_tokenizer	335
Tidying Up Input Text	337
Tokenizing HTML	337
Tidying Up Punctuation	338
20. Normalizing Tokens.....	341
In That Case	341

You Have an Accent	342
Retaining Meaning	343
Living in a Unicode World	346
Unicode Case Folding	347
Unicode Character Folding	349
Sorting and Collations	350
Case-Insensitive Sorting	351
Differences Between Languages	353
Unicode Collation Algorithm	353
Unicode Sorting	354
Specifying a Language	355
Customizing Collations	358
21. Reducing Words to Their Root Form.	359
Algorithmic Stemmers	360
Using an Algorithmic Stemmer	361
Dictionary Stemmers	363
Hunspell Stemmer	364
Installing a Dictionary	365
Per-Language Settings	365
Creating a Hunspell Token Filter	366
Hunspell Dictionary Format	367
Choosing a Stemmer	369
Stemmer Performance	370
Stemmer Quality	370
Stemmer Degree	370
Making a Choice	371
Controlling Stemming	371
Preventing Stemming	371
Customizing Stemming	372
Stemming in situ	373
Is Stemming in situ a Good Idea	375
22. Stopwords: Performance Versus Precision.	377
Pros and Cons of Stopwords	378
Using Stopwords	379
Stopwords and the Standard Analyzer	379
Maintaining Positions	380
Specifying Stopwords	380
Using the stop Token Filter	381
Updating Stopwords	383
Stopwords and Performance	383

and Operator	383
minimum_should_match	384
Divide and Conquer	385
Controlling Precision	386
Only High-Frequency Terms	387
More Control with Common Terms	388
Stopwords and Phrase Queries	388
Positions Data	389
Index Options	389
Stopwords	390
common_grams Token Filter	391
At Index Time	392
Unigram Queries	393
Bigram Phrase Queries	393
Two-Word Phrases	394
Stopwords and Relevance	394
23. Synonyms.....	395
Using Synonyms	396
Formatting Synonyms	397
Expand or contract	398
Simple Expansion	398
Simple Contraction	399
Genre Expansion	400
Synonyms and The Analysis Chain	401
Case-Sensitive Synonyms	401
Multiword Synonyms and Phrase Queries	402
Use Simple Contraction for Phrase Queries	404
Synonyms and the query_string Query	405
Symbol Synonyms	405
24. Typoes and Mispelings.....	409
Fuzziness	409
Fuzzy Query	410
Improving Performance	411
Fuzzy match Query	412
Scoring Fuzziness	413
Phonetic Matching	413

Part IV. Aggregations

25. High-Level Concepts.....	419
Buckets	420
Metrics	420
Combining the Two	420
26. Aggregation Test-Drive.....	423
Adding a Metric to the Mix	426
Buckets Inside Buckets	427
One Final Modification	429
27. Building Bar Charts.....	433
28. Looking at Time.....	437
Returning Empty Buckets	439
Extended Example	441
The Sky's the Limit	443
29. Scoping Aggregations.....	445
30. Filtering Queries and Aggregations.....	449
Filtered Query	449
Filter Bucket	450
Post Filter	451
Recap	452
31. Sorting Multivalued Buckets.....	453
Intrinsic Sorts	453
Sorting by a Metric	454
Sorting Based on “Deep” Metrics	455
32. Approximate Aggregations.....	457
Finding Distinct Counts	458
Understanding the Trade-offs	460
Optimizing for Speed	461
Calculating Percentiles	462
Percentile Metric	464
Percentile Ranks	467
Understanding the Trade-offs	469
33. Significant Terms.....	471
significant_terms Demo	472
Recommending Based on Popularity	474

Recommending Based on Statistics	478
34. Controlling Memory Use and Latency.....	481
Fielddata	481
Aggregations and Analysis	483
High-Cardinality Memory Implications	486
Limiting Memory Usage	487
Fielddata Size	488
Monitoring fielddata	489
Circuit Breaker	490
Fielddata Filtering	491
Doc Values	493
Enabling Doc Values	494
Preloading Fielddata	494
Eagerly Loading Fielddata	495
Global Ordinals	496
Index Warmers	498
Preventing Combinatorial Explosions	500
Depth-First Versus Breadth-First	502
35. Closing Thoughts.....	507

Part V. Geolocation

36. Geo-Points.....	511
Lat/Lon Formats	511
Filtering by Geo-Point	512
geo_bounding_box Filter	513
Optimizing Bounding Boxes	514
geo_distance Filter	515
Faster Geo-Distance Calculations	516
geo_distance_range Filter	517
Caching geo-filters	517
Reducing Memory Usage	519
Sorting by Distance	520
Scoring by Distance	522
37. Geohashes.....	523
Mapping Geohashes	524
geohash_cell Filter	525

38. Geo-aggregations.....	527
geo_distance Aggregation	527
geohash_grid Aggregation	530
geo_bounds Aggregation	532
39. Geo-shapes.....	535
Mapping geo-shapes	536
precision	536
distance_error_pct	537
Indexing geo-shapes	537
Querying geo-shapes	538
Querying with Indexed Shapes	540
Geo-shape Filters and Caching	541

Part VI. Modeling Your Data

40. Handling Relationships.....	545
Application-side Joins	546
Denormalizing Your Data	548
Field Collapsing	549
Denormalization and Concurrency	552
Renaming Files and Directories	555
Solving Concurrency Issues	555
Global Locking	556
Document Locking	557
Tree Locking	558
41. Nested Objects.....	561
Nested Object Mapping	563
Querying a Nested Object	564
Sorting by Nested Fields	565
Nested Aggregations	567
reverse_nested Aggregation	568
When to Use Nested Objects	570
42. Parent-Child Relationship.....	571
Parent-Child Mapping	572
Indexing Parents and Children	572
Finding Parents by Their Children	573
min_children and max_children	575
Finding Children by Their Parents	575

Children Aggregation	576
Grandparents and Grandchildren	577
Practical Considerations	579
Memory Use	579
Global Ordinals and Latency	580
Multigenerations and Concluding Thoughts	580
43. Designing for Scale.....	583
The Unit of Scale	583
Shard Overallocation	585
Kagillion Shards	586
Capacity Planning	587
Replica Shards	588
Balancing Load with Replicas	589
Multiple Indices	590
Time-Based Data	592
Index per Time Frame	592
Index Templates	593
Retiring Data	594
Migrate Old Indices	595
Optimize Indices	595
Closing Old Indices	596
Archiving Old Indices	596
User-Based Data	597
Shared Index	597
Faking Index per User with Aliases	600
One Big User	601
Scale Is Not Infinite	602

Part VII. Administration, Monitoring, and Deployment

44. Monitoring.....	607
Marvel for Monitoring	607
Cluster Health	608
Drilling Deeper: Finding Problematic Indices	609
Blocking for Status Changes	611
Monitoring Individual Nodes	612
indices Section	613
OS and Process Sections	616
JVM Section	617
Threadpool Section	620

FS and Network Sections	622
Circuit Breaker	622
Cluster Stats	623
Index Stats	623
Pending Tasks	624
cat API	626
45. Production Deployment.....	631
Hardware	631
Memory	631
CPUs	632
Disks	632
Network	633
General Considerations	633
Java Virtual Machine	634
Transport Client Versus Node Client	634
Configuration Management	635
Important Configuration Changes	635
Assign Names	636
Paths	636
Minimum Master Nodes	637
Recovery Settings	638
Prefer Unicast over Multicast	639
Don't Touch These Settings!	640
Garbage Collector	640
Threadpools	641
Heap: Sizing and Swapping	641
Give Half Your Memory to Lucene	642
Don't Cross 32 GB!	642
Swapping Is the Death of Performance	644
File Descriptors and MMap	645
Revisit This List Before Production	646
46. Post-Deployment.....	647
Changing Settings Dynamically	647
Logging	648
Slowlog	648
Indexing Performance Tips	649
Test Performance Scientifically	650
Using and Sizing Bulk Requests	650
Storage	651
Segments and Merging	651

Other	653
Rolling Restarts	654
Backing Up Your Cluster	655
Creating the Repository	655
Snapshotting All Open Indices	656
Snapshotting Particular Indices	657
Listing Information About Snapshots	657
Deleting Snapshots	658
Monitoring Snapshot Progress	658
Canceling a Snapshot	661
Restoring from a Snapshot	661
Monitoring Restore Operations	662
Canceling a Restore	663
Clusters Are Living, Breathing Creatures	664
Index.....	665

You Know, for Search...

Elasticsearch is an open-source search engine built on top of **Apache Lucene™**, a full-text search-engine library. Lucene is arguably the most advanced, high-performance, and fully featured search engine library in existence today—both open source and proprietary.

But Lucene is just a library. To leverage its power, you need to work in Java and to integrate Lucene directly with your application. Worse, you will likely require a degree in information retrieval to understand how it works. Lucene is *very* complex.

Elasticsearch is also written in Java and uses Lucene internally for all of its indexing and searching, but it aims to make full-text search easy by hiding the complexities of Lucene behind a simple, coherent, RESTful API.

However, Elasticsearch is much more than just Lucene and much more than “just” full-text search. It can also be described as follows:

- A distributed real-time document store where *every field* is indexed and searchable
- A distributed search engine with real-time analytics
- Capable of scaling to hundreds of servers and petabytes of structured and unstructured data

And it packages up all this functionality into a standalone server that your application can talk to via a simple RESTful API, using a web client from your favorite programming language, or even from the command line.

It is easy to get started with Elasticsearch. It ships with sensible defaults and hides complicated search theory away from beginners. It *just works*, right out of the box. With minimal understanding, you can soon become productive.

Elasticsearch can be downloaded, used, and modified free of charge. It is available under the [Apache 2 license](#), one of the most flexible open source licenses available.

As your knowledge grows, you can leverage more of Elasticsearch's advanced features. The entire engine is configurable and flexible. Pick and choose from the advanced features to tailor Elasticsearch to your problem domain.

The Mists of Time

Many years ago, a newly married unemployed developer called Shay Banon followed his wife to London, where she was studying to be a chef. While looking for gainful employment, he started playing with an early version of Lucene, with the intent of building his wife a recipe search engine.

Working directly with Lucene can be tricky, so Shay started work on an abstraction layer to make it easier for Java programmers to add search to their applications. He released this as his first open source project, called Compass.

Later Shay took a job working in a high-performance, distributed environment with in-memory data grids. The need for a high-performance, real-time, distributed search engine was obvious, and he decided to rewrite the Compass libraries as a standalone server called Elasticsearch.

The first public release came out in February 2010. Since then, Elasticsearch has become one of the most popular projects on GitHub with commits from over 300 contributors. A company has formed around Elasticsearch to provide commercial support and to develop new features, but Elasticsearch is, and forever will be, open source and available to all.

Shay's wife is still waiting for the recipe search...

Installing Elasticsearch

The easiest way to understand what Elasticsearch can do for you is to play with it, so let's get started!

The only requirement for installing Elasticsearch is a recent version of Java. Preferably, you should install the latest version of the official Java from www.java.com.

You can download the latest version of Elasticsearch from elasticsearch.org/download.

```
curl -L -O http://download.elasticsearch.org/PATH/TO/VERSION.zip ❶  
unzip elasticsearch-$VERSION.zip  
cd elasticsearch-$VERSION
```

- ❶ Fill in the URL for the latest version available on elasticsearch.org/download.



When installing Elasticsearch in production, you can use the method described previously, or the Debian or RPM packages provided on the [downloads page](#). You can also use the officially supported [Puppet module](#) or [Chef cookbook](#).

Installing Marvel

Marvel is a management and monitoring tool for Elasticsearch, which is free for development use. It comes with an interactive console called Sense, which makes it easy to talk to Elasticsearch directly from your browser.

Many of the code examples in the online version of this book include a View in Sense link. When clicked, it will open up a working example of the code in the Sense console. You do not have to install Marvel, but it will make this book much more interactive by allowing you to experiment with the code samples on your local Elasticsearch cluster.

Marvel is available as a plug-in. To download and install it, run this command in the Elasticsearch directory:

```
./bin/plugin -i elasticsearch/marvel/latest
```

You probably don't want Marvel to monitor your local cluster, so you can disable data collection with this command:

```
echo 'marvel.agent.enabled: false' >> ./config/elasticsearch.yml
```

Running Elasticsearch

Elasticsearch is now ready to run. You can start it up in the foreground with this:

```
./bin/elasticsearch
```

Add `-d` if you want to run it in the background as a daemon.

Test it out by opening another terminal window and running the following:

```
curl 'http://localhost:9200/?pretty'
```

You should see a response like this:

```
{
  "status": 200,
  "name": "Shrunken Bones",
  "version": {
    "number": "1.4.0",
    "lucene_version": "4.10"
  },
  "tagline": "You Know, for Search"
}
```

This means that your Elasticsearch *cluster* is up and running, and we can start experimenting with it.



A *node* is a running instance of Elasticsearch. A *cluster* is a group of nodes with the same `cluster.name` that are working together to share data and to provide failover and scale, although a single node can form a cluster all by itself.

You should change the default `cluster.name` to something appropriate to you, like your own name, to stop your nodes from trying to join another cluster on the same network with the same name!

You can do this by editing the `elasticsearch.yml` file in the `config/` directory and then restarting Elasticsearch. When Elasticsearch is running in the foreground, you can stop it by pressing Ctrl-C; otherwise, you can shut it down with the shutdown API:

```
curl -XPOST 'http://localhost:9200/_shutdown'
```

Viewing Marvel and Sense

If you installed the **Marvel** management and monitoring tool, you can view it in a web browser by visiting http://localhost:9200/_plugin/marvel/.

You can reach the *Sense* developer console either by clicking the “Marvel dashboards” drop-down in Marvel, or by visiting http://localhost:9200/_plugin/marvel/sense/.

Talking to Elasticsearch

How you talk to Elasticsearch depends on whether you are using Java.

Java API

If you are using Java, Elasticsearch comes with two built-in clients that you can use in your code:

Node client

The node client joins a local cluster as a *non data node*. In other words, it doesn’t hold any data itself, but it knows what data lives on which node in the cluster, and can forward requests directly to the correct node.

Transport client

The lighter-weight transport client can be used to send requests to a remote cluster. It doesn’t join the cluster itself, but simply forwards requests to a node in the cluster.

Both Java clients talk to the cluster over port 9300, using the native Elasticsearch *transport* protocol. The nodes in the cluster also communicate with each other over port 9300. If this port is not open, your nodes will not be able to form a cluster.



The Java client must be from the same version of Elasticsearch as the nodes; otherwise, they may not be able to understand each other.

More information about the Java clients can be found in the Java API section of the [Guide](#).

RESTful API with JSON over HTTP

All other languages can communicate with Elasticsearch over port 9200 using a RESTful API, accessible with your favorite web client. In fact, as you have seen, you can even talk to Elasticsearch from the command line by using the `curl` command.



Elasticsearch provides official clients for several languages—Groovy, JavaScript, .NET, PHP, Perl, Python, and Ruby—and there are numerous community-provided clients and integrations, all of which can be found in the [Guide](#).

A request to Elasticsearch consists of the same parts as any HTTP request:

```
curl -X<VERB> '<PROTOCOL>://<HOST>/<PATH>?<QUERY_STRING>' -d '<BODY>'
```

The parts marked with `< >` above are:

VERB

The appropriate HTTP *method* or *verb*: GET, POST, PUT, HEAD, or DELETE.

PROTOCOL

Either `http` or `https` (if you have an `https` proxy in front of Elasticsearch.)

HOST

The hostname of any node in your Elasticsearch cluster, or `localhost` for a node on your local machine.

PORT

The port running the Elasticsearch HTTP service, which defaults to 9200.

QUERY_STRING

Any optional query-string parameters (for example `?pretty` will *pretty-print* the JSON response to make it easier to read.)

BODY

A JSON-encoded request body (if the request needs one.)

For instance, to count the number of documents in the cluster, we could use this:

```
curl -XGET 'http://localhost:9200/_count?pretty' -d '{
  "query": {
    "match_all": {}
  }
}'
```

Elasticsearch returns an HTTP status code like 200 OK and (except for HEAD requests) a JSON-encoded response body. The preceding `curl` request would respond with a JSON body like the following:

```
{
  "count" : 0,
  "_shards" : {
    "total" : 5,
    "successful" : 5,
    "failed" : 0
  }
}
```

We don't see the HTTP headers in the response because we didn't ask `curl` to display them. To see the headers, use the `curl` command with the `-i` switch:

```
curl -i -XGET 'localhost:9200/'
```

For the rest of the book, we will show these `curl` examples using a shorthand format that leaves out all the bits that are the same in every request, like the hostname and port, and the `curl` command itself. Instead of showing a full request like

```
curl -XGET 'localhost:9200/_count?pretty' -d '{
  "query": {
    "match_all": {}
  }
}'
```

we will show it in this shorthand format:

```
GET /_count
{
  "query": {
    "match_all": {}
  }
}
```

In fact, this is the same format that is used by the Sense console that we installed with **Marvel**. If in the online version of this book, you can open and run this code example in Sense by clicking the View in Sense link above.

Document Oriented

Objects in an application are seldom just a simple list of keys and values. More often than not, they are complex data structures that may contain dates, geo locations, other objects, or arrays of values.

Sooner or later you're going to want to store these objects in a database. Trying to do this with the rows and columns of a relational database is the equivalent of trying to squeeze your rich, expressive objects into a very big spreadsheet: you have to flatten the object to fit the table schema—usually one field per column—and then have to reconstruct it every time you retrieve it.

Elasticsearch is *document oriented*, meaning that it stores entire objects or *documents*. It not only stores them, but also *indexes* the contents of each document in order to make them searchable. In Elasticsearch, you index, search, sort, and filter documents—not rows of columnar data. This is a fundamentally different way of thinking about data and is one of the reasons Elasticsearch can perform complex full-text search.

JSON

Elasticsearch uses JavaScript Object Notation, or **JSON**, as the serialization format for documents. JSON serialization is supported by most programming languages, and has become the standard format used by the NoSQL movement. It is simple, concise, and easy to read.

Consider this JSON document, which represents a user object:

```
{
  "email":      "john@smith.com",
  "first_name": "John",
  "last_name":  "Smith",
  "info": {
    "bio":      "Eco-warrior and defender of the weak",
    "age":      25,
    "interests": [ "dolphins", "whales" ]
  },
  "join_date": "2014/05/01"
}
```

Although the original user object was complex, the structure and meaning of the object has been retained in the JSON version. Converting an object to JSON for indexing in Elasticsearch is much simpler than the equivalent process for a flat table structure.



Almost all languages have modules that will convert arbitrary data structures or objects into JSON for you, but the details are specific to each language. Look for modules that handle JSON *serialization* or *marshalling*. **The official Elasticsearch clients** all handle conversion to and from JSON for you automatically.

Finding Your Feet

To give you a feel for what is possible in Elasticsearch and how easy it is to use, let's start by walking through a simple tutorial that covers basic concepts such as indexing, search, and aggregations.

We'll introduce some new terminology and basic concepts along the way, but it is OK if you don't understand everything immediately. We'll cover all the concepts introduced here in *much* greater depth throughout the rest of the book.

So, sit back and enjoy a whirlwind tour of what Elasticsearch is capable of.

Let's Build an Employee Directory

We happen to work for *Megacorp*, and as part of HR's new "*We love our drones!*" initiative, we have been tasked with creating an employee directory. The directory is supposed to foster employer empathy and real-time, synergistic, dynamic collaboration, so it has a few business requirements:

- Enable data to contain multi value tags, numbers, and full text.
- Retrieve the full details of any employee.
- Allow structured search, such as finding employees over the age of 30.
- Allow simple full-text search and more-complex *phrase* searches.
- Return highlighted search *snippets* from the text in the matching documents.
- Enable management to build analytic dashboards over the data.

Indexing Employee Documents

The first order of business is storing employee data. This will take the form of an *employee document*: a single document represents a single employee. The act of storing data in Elasticsearch is called *indexing*, but before we can index a document, we need to decide *where* to store it.

In Elasticsearch, a document belongs to a *type*, and those types live inside an *index*. You can draw some (rough) parallels to a traditional relational database:

Relational DB ⇒ Databases ⇒ Tables ⇒ Rows ⇒ Columns
Elasticsearch ⇒ Indices ⇒ Types ⇒ Documents ⇒ Fields

An Elasticsearch cluster can contain multiple *indices* (databases), which in turn contain multiple *types* (tables). These types hold multiple *documents* (rows), and each document has multiple *fields* (columns).

Index Versus Index Versus Index

You may already have noticed that the word *index* is overloaded with several meanings in the context of Elasticsearch. A little clarification is necessary:

Index (noun)

As explained previously, an *index* is like a *database* in a traditional relational database. It is the place to store related documents. The plural of *index* is *indices* or *indexes*.

Index (verb)

To *index a document* is to store a document in an *index (noun)* so that it can be retrieved and queried. It is much like the INSERT keyword in SQL except that, if the document already exists, the new document would replace the old.

Inverted index

Relational databases add an *index*, such as a B-tree index, to specific columns in order to improve the speed of data retrieval. Elasticsearch and Lucene use a structure called an *inverted index* for exactly the same purpose.

By default, every field in a document is *indexed* (has an inverted index) and thus is searchable. A field without an inverted index is not searchable. We discuss inverted indexes in more detail in **“Inverted Index” on page 81**.

So for our employee directory, we are going to do the following:

- Index a *document* per employee, which contains all the details of a single employee.
- Each document will be of *type* employee.
- That type will live in the megacorp *index*.
- That index will reside within our Elasticsearch cluster.

In practice, this is easy (even though it looks like a lot of steps). We can perform all of those actions in a single command:

```
PUT /megacorp/employee/1
{
  "first_name" : "John",
  "last_name"  : "Smith",
```

```

    "age" :      25,
    "about" :    "I love to go rock climbing",
    "interests": [ "sports", "music" ]
  }

```

Notice that the path `/megacorp/employee/1` contains three pieces of information:

`megacorp`

The index name

`employee`

The type name

`1`

The ID of this particular employee

The request body—the JSON document—contains all the information about this employee. His name is John Smith, he's 25, and enjoys rock climbing.

Simple! There was no need to perform any administrative tasks first, like creating an index or specifying the type of data that each field contains. We could just index a document directly. Elasticsearch ships with defaults for everything, so all the necessary administration tasks were taken care of in the background, using default values.

Before moving on, let's add a few more employees to the directory:

```

PUT /megacorp/employee/2
{
  "first_name" : "Jane",
  "last_name"  : "Smith",
  "age" :      32,
  "about" :    "I like to collect rock albums",
  "interests": [ "music" ]
}

PUT /megacorp/employee/3
{
  "first_name" : "Douglas",
  "last_name"  : "Fir",
  "age" :      35,
  "about" :    "I like to build cabinets",
  "interests": [ "forestry" ]
}

```

Retrieving a Document

Now that we have some data stored in Elasticsearch, we can get to work on the business requirements for this application. The first requirement is the ability to retrieve individual employee data.

This is easy in Elasticsearch. We simply execute an HTTP GET request and specify the *address* of the document—the index, type, and ID. Using those three pieces of information, we can return the original JSON document:

```
GET /megacorp/employee/1
```

And the response contains some metadata about the document, and John Smith's original JSON document as the `_source` field:

```
{
  "_index" : "megacorp",
  "_type" : "employee",
  "_id" : "1",
  "_version" : 1,
  "found" : true,
  "_source" : {
    "first_name" : "John",
    "last_name" : "Smith",
    "age" : 25,
    "about" : "I love to go rock climbing",
    "interests": [ "sports", "music" ]
  }
}
```



In the same way that we changed the HTTP verb from PUT to GET in order to retrieve the document, we could use the DELETE verb to delete the document, and the HEAD verb to check whether the document exists. To replace an existing document with an updated version, we just PUT it again.

Search Lite

A GET is fairly simple—you get back the document that you ask for. Let's try something a little more advanced, like a simple search!

The first search we will try is the simplest search possible. We will search for all employees, with this request:

```
GET /megacorp/employee/_search
```

You can see that we're still using index `megacorp` and type `employee`, but instead of specifying a document ID, we now use the `_search` endpoint. The response includes all three of our documents in the `hits` array. By default, a search will return the top 10 results.

```
{
  "took": 6,
  "timed_out": false,
  "_shards": { ... },
  "hits": {
```

```

"total":      3,
"max_score":  1,
"hits": [
  {
    "_index":      "megacorp",
    "_type":      "employee",
    "_id":        "3",
    "_score":      1,
    "_source": {
      "first_name": "Douglas",
      "last_name":  "Fir",
      "age":        35,
      "about":      "I like to build cabinets",
      "interests":  [ "forestry" ]
    }
  },
  {
    "_index":      "megacorp",
    "_type":      "employee",
    "_id":        "1",
    "_score":      1,
    "_source": {
      "first_name": "John",
      "last_name":  "Smith",
      "age":        25,
      "about":      "I love to go rock climbing",
      "interests":  [ "sports", "music" ]
    }
  },
  {
    "_index":      "megacorp",
    "_type":      "employee",
    "_id":        "2",
    "_score":      1,
    "_source": {
      "first_name": "Jane",
      "last_name":  "Smith",
      "age":        32,
      "about":      "I like to collect rock albums",
      "interests":  [ "music" ]
    }
  }
]
}

```



The response not only tells us which documents matched, but also includes the whole document itself: all the information that we need in order to display the search results to the user.

Next, let's try searching for employees who have "Smith" in their last name. To do this, we'll use a *lightweight* search method that is easy to use from the command line. This method is often referred to as a *query-string* search, since we pass the search as a URL query-string parameter:

```
GET /megacorp/employee/_search?q=last_name:Smith
```

We use the same `_search` endpoint in the path, and we add the query itself in the `q=` parameter. The results that come back show all Smiths:

```
{
  ...
  "hits": {
    "total": 2,
    "max_score": 0.30685282,
    "hits": [
      {
        ...
        "_source": {
          "first_name": "John",
          "last_name": "Smith",
          "age": 25,
          "about": "I love to go rock climbing",
          "interests": [ "sports", "music" ]
        }
      },
      {
        ...
        "_source": {
          "first_name": "Jane",
          "last_name": "Smith",
          "age": 32,
          "about": "I like to collect rock albums",
          "interests": [ "music" ]
        }
      }
    ]
  }
}
```

Search with Query DSL

Query-string search is handy for ad hoc searches from the command line, but it has its limitations (see [“Search Lite” on page 76](#)). Elasticsearch provides a rich, flexible, query language called the *query DSL*, which allows us to build much more complicated, robust queries.

The *domain-specific language* (DSL) is specified using a JSON request body. We can represent the previous search for all Smiths like so:

```
GET /megacorp/employee/_search
{
  "query" : {
    "match" : {
      "last_name" : "Smith"
    }
  }
}
```

This will return the same results as the previous query. You can see that a number of things have changed. For one, we are no longer using *query-string* parameters, but instead a request body. This request body is built with JSON, and uses a *match* query (one of several types of queries, which we will learn about later).

More-Complicated Searches

Let's make the search a little more complicated. We still want to find all employees with a last name of Smith, but we want only employees who are older than 30. Our query will change a little to accommodate a *filter*, which allows us to execute structured searches efficiently:

```
GET /megacorp/employee/_search
{
  "query" : {
    "filtered" : {
      "filter" : {
        "range" : {
          "age" : { "gt" : 30 } ❶
        }
      },
      "query" : {
        "match" : {
          "last_name" : "smith" ❷
        }
      }
    }
  }
}
```

❶ This portion of the query is a range *filter*, which will find all ages older than 30—gt stands for *greater than*.

❷ This portion of the query is the same *match query* that we used before.

Don't worry about the syntax too much for now; we will cover it in great detail later. Just recognize that we've added a *filter* that performs a range search, and reused the same *match query* as before. Now our results show only one employee who happens to be 32 and is named Jane Smith:

```

{
  ...
  "hits": {
    "total": 1,
    "max_score": 0.30685282,
    "hits": [
      {
        ...
        "_source": {
          "first_name": "Jane",
          "last_name": "Smith",
          "age": 32,
          "about": "I like to collect rock albums",
          "interests": [ "music" ]
        }
      }
    ]
  }
}

```

Full-Text Search

The searches so far have been simple: single names, filtered by age. Let's try a more advanced, full-text search—a task that traditional databases would really struggle with.

We are going to search for all employees who enjoy rock climbing:

```

GET /megacorp/employee/_search
{
  "query" : {
    "match" : {
      "about" : "rock climbing"
    }
  }
}

```

You can see that we use the same `match` query as before to search the `about` field for “rock climbing.” We get back two matching documents:

```

{
  ...
  "hits": {
    "total": 2,
    "max_score": 0.16273327,
    "hits": [
      {
        ...
        "_score": 0.16273327, ❶
        "_source": {
          "first_name": "John",
          "last_name": "Smith",

```

```

        "age":      25,
        "about":    "I love to go rock climbing",
        "interests": [ "sports", "music" ]
    },
    {
        ...
        "_score":    0.016878016, ❶
        "_source": {
            "first_name": "Jane",
            "last_name":  "Smith",
            "age":        32,
            "about":      "I like to collect rock albums",
            "interests": [ "music" ]
        }
    }
]
}

```

❶ The relevance scores

By default, Elasticsearch sorts matching results by their relevance score, that is, by how well each document matches the query. The first and highest-scoring result is obvious: John Smith’s about field clearly says “rock climbing” in it.

But why did Jane Smith come back as a result? The reason her document was returned is because the word “rock” was mentioned in her about field. Because only “rock” was mentioned, and not “climbing,” her `_score` is lower than John’s.

This is a good example of how Elasticsearch can search *within* full-text fields and return the most relevant results first. This concept of *relevance* is important to Elasticsearch, and is a concept that is completely foreign to traditional relational databases, in which a record either matches or it doesn’t.

Phrase Search

Finding individual words in a field is all well and good, but sometimes you want to match exact sequences of words or *phrases*. For instance, we could perform a query that will match only employee records that contain both “rock” *and* “climbing” *and* that display the words are next to each other in the phrase “rock climbing.”

To do this, we use a slight variation of the `match` query called the `match_phrase` query:

```

GET /megacorp/employee/_search
{
  "query" : {
    "match_phrase" : {

```

```

        "about" : "rock climbing"
    }
}

```

This, to no surprise, returns only John Smith's document:

```

{
  ...
  "hits": {
    "total":      1,
    "max_score":  0.23013961,
    "hits": [
      {
        ...
        "_score":      0.23013961,
        "_source": {
          "first_name": "John",
          "last_name":  "Smith",
          "age":        25,
          "about":      "I love to go rock climbing",
          "interests": [ "sports", "music" ]
        }
      }
    ]
  }
}

```

Highlighting Our Searches

Many applications like to *highlight* snippets of text from each search result so the user can see *why* the document matched the query. Retrieving highlighted fragments is easy in Elasticsearch.

Let's rerun our previous query, but add a new `highlight` parameter:

```

GET /megacorp/employee/_search
{
  "query" : {
    "match_phrase" : {
      "about" : "rock climbing"
    }
  },
  "highlight": {
    "fields" : {
      "about" : {}
    }
  }
}

```

When we run this query, the same hit is returned as before, but now we get a new section in the response called `highlight`. This contains a snippet of text from the `about` field with the matching words wrapped in `` HTML tags:

```
{
  ...
  "hits": {
    "total": 1,
    "max_score": 0.23013961,
    "hits": [
      {
        ...
        "_score": 0.23013961,
        "_source": {
          "first_name": "John",
          "last_name": "Smith",
          "age": 25,
          "about": "I love to go rock climbing",
          "interests": [ "sports", "music" ]
        },
        "highlight": {
          "about": [
            "I love to go <em>rock</em> <em>climbing</em>" ❶
          ]
        }
      }
    ]
  }
}
```

❶ The highlighted fragment from the original text

You can read more about the highlighting of search snippets in the [highlighting reference documentation](#).

Analytics

Finally, we come to our last business requirement: allow managers to run analytics over the employee directory. Elasticsearch has functionality called *aggregations*, which allow you to generate sophisticated analytics over your data. It is similar to `GROUP BY` in SQL, but much more powerful.

For example, let's find the most popular interests enjoyed by our employees:

```
GET /megacorp/employee/_search
{
  "aggs": {
    "all_interests": {
      "terms": { "field": "interests" }
    }
  }
}
```

```
}
}
```

Ignore the syntax for now and just look at the results:

```
{
  ...
  "hits": { ... },
  "aggregations": {
    "all_interests": {
      "buckets": [
        {
          "key": "music",
          "doc_count": 2
        },
        {
          "key": "forestry",
          "doc_count": 1
        },
        {
          "key": "sports",
          "doc_count": 1
        }
      ]
    }
  }
}
```

We can see that two employees are interested in music, one in forestry, and one in sports. These aggregations are not precalculated; they are generated on the fly from the documents that match the current query. If we want to know the popular interests of people called Smith, we can just add the appropriate query into the mix:

```
GET /megacorp/employee/_search
{
  "query": {
    "match": {
      "last_name": "smith"
    }
  },
  "aggs": {
    "all_interests": {
      "terms": {
        "field": "interests"
      }
    }
  }
}
```

The `all_interests` aggregation has changed to include only documents matching our query:

```

...
"all_interests": {
  "buckets": [
    {
      "key": "music",
      "doc_count": 2
    },
    {
      "key": "sports",
      "doc_count": 1
    }
  ]
}

```

Aggregations allow hierarchical rollups too. For example, let's find the average age of employees who share a particular interest:

```

GET /megacorp/employee/_search
{
  "aggs" : {
    "all_interests" : {
      "terms" : { "field" : "interests" },
      "aggs" : {
        "avg_age" : {
          "avg" : { "field" : "age" }
        }
      }
    }
  }
}

```

The aggregations that we get back are a bit more complicated, but still fairly easy to understand:

```

...
"all_interests": {
  "buckets": [
    {
      "key": "music",
      "doc_count": 2,
      "avg_age": {
        "value": 28.5
      }
    },
    {
      "key": "forestry",
      "doc_count": 1,
      "avg_age": {
        "value": 35
      }
    },
    {
      "key": "sports",

```



```

        "doc_count": 1,
        "avg_age": {
          "value": 25
        }
      ]
    }
  }
}

```

The output is basically an enriched version of the first aggregation we ran. We still have a list of interests and their counts, but now each interest has an additional `avg_age`, which shows the average age for all employees having that interest.

Even if you don't understand the syntax yet, you can easily see how complex aggregations and groupings can be accomplished using this feature. The sky is the limit as to what kind of data you can extract!

Tutorial Conclusion

Hopefully, this little tutorial was a good demonstration about what is possible in Elasticsearch. It is really just scratching the surface, and many features—such as suggestions, geolocation, percolation, fuzzy and partial matching—were omitted to keep the tutorial short. But it did highlight just how easy it is to start building advanced search functionality. No configuration was needed—just add data and start searching!

It's likely that the syntax left you confused in places, and you may have questions about how to tweak and tune various aspects. That's fine! The rest of the book dives into each of these issues in detail, giving you a solid understanding of how Elasticsearch works.

Distributed Nature

At the beginning of this chapter, we said that Elasticsearch can scale out to hundreds (or even thousands) of servers and handle petabytes of data. While our tutorial gave examples of how to use Elasticsearch, it didn't touch on the mechanics at all. Elasticsearch is distributed by nature, and it is designed to hide the complexity that comes with being distributed.

The distributed aspect of Elasticsearch is largely transparent. Nothing in the tutorial required you to know about distributed systems, sharding, cluster discovery, or dozens of other distributed concepts. It happily ran the tutorial on a single node living inside your laptop, but if you were to run the tutorial on a cluster containing 100 nodes, everything would work in exactly the same way.

Elasticsearch tries hard to hide the complexity of distributed systems. Here are some of the operations happening automatically under the hood:

- Partitioning your documents into different containers or *shards*, which can be stored on a single node or on multiple nodes
- Balancing these shards across the nodes in your cluster to spread the indexing and search load
- Duplicating each shard to provide redundant copies of your data, to prevent data loss in case of hardware failure
- Routing requests from any node in the cluster to the nodes that hold the data you're interested in
- Seamlessly integrating new nodes as your cluster grows or redistributing shards to recover from node loss

As you read through this book, you'll encounter supplemental chapters about the distributed nature of Elasticsearch. These chapters will teach you about how the cluster scales and deals with failover ([Chapter 2](#)), handles document storage ([Chapter 4](#)), executes distributed search ([Chapter 9](#)), and what a shard is and how it works ([Chapter 11](#)).

These chapters are not required reading—you can use Elasticsearch without understanding these internals—but they will provide insight that will make your knowledge of Elasticsearch more complete. Feel free to skim them and revisit at a later point when you need a more complete understanding.

Next Steps

By now you should have a taste of what you can do with Elasticsearch, and how easy it is to get started. Elasticsearch tries hard to work out of the box with minimal knowledge and configuration. The best way to learn Elasticsearch is by jumping in: just start indexing and searching!

However, the more you know about Elasticsearch, the more productive you can become. The more you can tell Elasticsearch about the domain-specific elements of your application, the more you can fine-tune the output.

The rest of this book will help you move from novice to expert. Each chapter explains the essentials, but also includes expert-level tips. If you're just getting started, these tips are probably not immediately relevant to you; Elasticsearch has sensible defaults and will generally do the right thing without any interference. You can always revisit these chapters later, when you are looking to improve performance by shaving off any wasted milliseconds.

Want to read more?

You can [buy this book](#) at **oreilly.com**
in print and ebook format.

Buy 2 books, get the 3rd FREE!

Use discount code: OPC10

All orders over \$29.95 qualify for **free shipping** within the US.

It's also available at your favorite book retailer,
including the iBookstore, the [Android Marketplace](#),
and [Amazon.com](#).



O'REILLY®

Spreading the knowledge of innovators

[oreilly.com](#)